

# AIF MuJoCo Robot

*Mathematical Reference*

---

Active Inference controller for a 3D MuJoCo robot.  
Beliefs, generative model, EFE, policy, and RxInfer factor graph.

# 1. State-Space Model

---

The robot operates in 3D Cartesian space ( $x, y, z$ ). The state is the position vector  $s = [s_x, s_y, s_z]$ . The dynamics are additive: control  $u$  shifts the state directly.

## Transition Model

Given state  $s$  and control  $u$ , the next state is:

$$s' = s + u$$

With per-dimension process noise  $Q = \text{diag}(q_x, q_y, q_z)$ , the stochastic transition becomes:

$$s'_i = s_i + u_i + \sigma_i, \quad \sigma_i \sim N(0, q_i)$$

where  $i \in \{x, y, z\}$ .

## Observation Model

Observations are noisy readings of the true state:

$$y_i \sim N(s_i, r_i)$$

where  $r_i$  is the observation noise variance for axis  $i$ . The observation noise can be a scalar (same for all axes) or per-axis  $[r_x, r_y, r_z]$ .

*Source: src/aif/generative\_model.jl*

## 2. Belief State

---

The agent maintains a Gaussian belief over its 3D position with a diagonal covariance matrix:

$$b(s) = N(\mu, \text{diag}(\sigma^2_x, \sigma^2_y, \sigma^2_z))$$

where  $\mu = [\mu_x, \mu_y, \mu_z]$  is the belief mean and  $\sigma^2_i$  is the marginal variance for axis i.

### Covariance Bounds

To prevent numerical instability:

$$1 \times 10^{-5} \leq \sigma^2_i \leq 2.0$$

All covariance updates are clamped to these bounds.

Source: `src/aif/beliefs.jl` | `COV_MIN = 1e-5, COV_MAX = 2.0`

## 3. Prediction Step

---

Before observing, the belief is propagated forward using the applied control (the actual scaled control, not the raw action):

### Mean Prediction

$$\mu'_{\cdot i} = \mu_i + \text{ctrl}_i$$

### Covariance Prediction

$$\sigma^2'_{\cdot i} = \text{clamp}(\sigma^2_i + q_i, \text{COV\_MIN}, \text{COV\_MAX})$$

Process noise  $q$  can be a scalar (applied equally to all axes) or a 3-vector [ $q_x, q_y, q_z$ ]. Default:  $q = 0.002$ .

*Note: The prediction uses the actual applied control ( $\text{ctrl} = \text{clamp}(\text{action} \times \text{scale})$ ), not the raw action, to match the true dynamics.*

Source: [src/aif/beliefs.jl](#) | `predict_belief!()`

## 4. Bayesian Update (Observation)

---

When an observation  $y$  arrives, the belief is updated using precision-weighted fusion (conjugate Gaussian update), independently per axis:

### Precision Fusion

$$\tau_{\text{prior}_i} = 1 / \sigma^2_i$$

$$\tau_{\text{lik}_i} = 1 / r_i$$

$$\tau_{\text{post}_i} = \tau_{\text{prior}_i} + \tau_{\text{lik}_i}$$

### Posterior Covariance

$$\sigma^2_{\text{post}_i} = \text{clamp}(1 / \tau_{\text{post}_i}, \text{COV\_MIN}, \text{COV\_MAX})$$

### Posterior Mean

$$\mu_{\text{post}_i} = (\tau_{\text{prior}_i} \cdot \mu_i + \tau_{\text{lik}_i} \cdot y_i) / \tau_{\text{post}_i}$$

This is the standard Kalman filter update for the diagonal case: the posterior mean is a precision-weighted average of the prior mean and the observation.

*Source: src/aif/beliefs.jl | update\_belief!()*

## 5. Expected Free Energy (EFE)

Policy selection minimizes Expected Free Energy. EFE decomposes into a pragmatic (goal-seeking) term and an epistemic (exploration) term:

$$\text{EFE}(a) = G_{\text{pragmatic}}(a) - G_{\text{epistemic}}$$

### Pragmatic Term (Goal-Seeking)

Measures expected squared distance to goal after applying action  $a$ . The predicted position uses the actual control scaling:

$$\text{pred}_i = \mu_i + \text{clamp}(a_i \times \text{ctrl\_scale}, -\text{ctrl\_lim}, \text{ctrl\_lim})$$

$$G_{\text{pragmatic}} = \gamma \cdot \sum_i w_i \cdot (\text{pred}_i - \text{goal}_i)^2$$

where  $\gamma$  is the pragmatic weight (default 1.5) and  $w_i$  are optional per-axis weights (default [1, 1, 1]).

### Epistemic Term (Exploration)

Encourages uncertainty reduction based on belief entropy:

$$G_{\text{epistemic}} = -\beta \cdot \sum_i \log(\sigma^2_i)$$

where  $\beta$  is the epistemic weight (default 0.02). Covariance is clamped to [1e-8, 100] for numerical stability in the log.

### Combined EFE

$$\text{EFE}(a) = \gamma \sum_i w_i (\text{pred}_i - \text{goal}_i)^2 + \beta \sum_i \log(\sigma^2_i)$$

Lower EFE = better action (closer to goal + reduces uncertainty).

Source: [src/aif/efe.jl](#)

## 6. Policy Selection

---

The agent selects the action that minimizes EFE over a discrete action set:

$$a^* = \operatorname{argmin}_{\{a \in A\}} \text{EFE}(a)$$

### Action Set

$A = 7 \times 7 \times 7 = 343$  actions. Each axis has 7 velocity levels:  $\{-3s, -2s, -s, 0, s, 2s, 3s\}$  with  $\text{step\_size } s = 0.04$ .

$$A = \{ [dx, dy, dz] : dx, dy, dz \in \{-0.12, -0.08, -0.04, 0, 0.04, 0.08, 0.12\} \}$$

The fine 7-level grid minimises quantization artifacts for smoother trajectories. Combined with EMA control smoothing, this produces clean, jitter-free curves.

*Source: src/aif/policy.jl*

## 7. Action to Control Mapping

---

The raw action  $a$  is scaled and clamped to produce the MuJoCo control signal:

```
ctrl_i = clamp(a_i * scale, -ctrl_lim, ctrl_lim)
```

Default:  $scale = 3.0$ ,  $ctrl\_lim = 1.2$ . This prevents overshooting while allowing sufficient movement per step.

### EMA Smoothing

Before applying the control, an Exponential Moving Average (EMA) filter smooths consecutive control signals:

```
ctrl_smooth = alpha * ctrl_new + (1 - alpha) * ctrl_prev
```

where  $\alpha$  is the smoothing weight (default 0.3). Lower  $\alpha$  produces smoother trajectories (more weight on previous control).  $\alpha = 1.0$  disables smoothing.

The control vector  $ctrl = [ctrl_x, ctrl_y, ctrl_z]$  is added to the current position in MuJoCo:  $target = pos + ctrl$ .

Source: [src/aif/action.jl](#)

## 8. RxInfer Factor Graph

An alternative inference backend expresses the same state-space model as a factor graph using RxInfer.jl. Each axis (x, y, z) runs an independent 1D linear-Gaussian model via reactive message passing.

### Per-Axis Factor Graph

```
x_prev ~ Normal(mean = m_prev, variance = v_prev)      [prior]
x      ~ Normal(mean = x_prev + u, variance = q)       [transition]
y      ~ Normal(mean = x, variance = r)                 [observation]
```

where  $m_{\text{prev}}$  and  $v_{\text{prev}}$  are the posterior parameters from the previous timestep,  $u$  is the applied control,  $q$  is process noise variance, and  $r$  is observation noise variance.

### Message Passing Schedule

At each timestep t:

- Forward message:  $x_{\text{prev}} \rightarrow x$  carries  $N(m_{\text{prev}} + u, v_{\text{prev}} + q)$
- Likelihood message:  $y \rightarrow x$  carries  $N(y, r)$
- Posterior:  $q(x) = N(\mu_{\text{post}}, \sigma^2_{\text{post}})$  via precision fusion

### Autoupdates (Prior Propagation)

After each inference cycle, the posterior becomes the next prior:

```
m_prev, v_prev → mean(q(x)), var(q(x))
```

This is implemented with RxInfer's `@autoupdates` macro and a single atomic `RecentSubject` stream to avoid synchronization issues.

### Equivalence to Analytic Update

For the linear-Gaussian model, belief propagation is exact. The RxInfer posterior matches the closed-form Kalman update to machine epsilon ( $\approx 2.2 \times 10^{-16}$ ). Both backends are interchangeable.

*Source: [src/aif/rxinfer\\_filter.jl](#)*

## 9. Utility Functions

---

### Normalize

Normalizes a probability vector in place:

$$p_i \rightarrow p_i / \sum_j p_j \quad (\text{if } \sum > 0, \text{ else uniform})$$

### Gaussian PDF

Univariate:

$$p(x | \mu, \sigma^2) = (1 / \sqrt{2\pi\sigma^2}) \cdot \exp(-(x-\mu)^2 / (2\sigma^2))$$

Multivariate (diagonal covariance):

$$p(x | \mu, \Sigma) = \prod_i p(x_i | \mu_i, \sigma^2_i)$$

### Gaussian Entropy

Entropy of a univariate Gaussian:

$$H = 0.5 \cdot \log(2\pi e \sigma^2)$$

For the diagonal belief state, the total entropy is the sum over axes.

### Softmax

Numerically stable softmax (subtract max for overflow protection):

$$\text{softmax}(x)_i = \exp(x_i - \max(x)) / \sum_j \exp(x_j - \max(x))$$

Source: [src/utils/math.jl](#)